



RV College of Engineering®

Mysore Road, RV Vidyaniketan Post, Bengaluru - 560059, Karnataka, India

Go, change the world

NeuroLearn - AI based cognitive learning styles and content generation

An Interdisciplinary Project Report (XX367P)

Submitted by,

Aniket R T

1RV23AI017

Alroy Deon Saldanha

1RV23AI012

Smriti J Abraham

1RV23ME106

Sudhanva Patil

1RV23ME113

Under the guidance of

Dr. Anupama Kumar

Assistant Professor

Dept. of AIML

RV College of Engineering®

In partial fulfillment of the requirements for the degree of
Bachelor of Engineering in respective departments

2025-26



RV College of Engineering®, Bengaluru
(Autonomous institution affiliated to VTU, Belagavi)



CERTIFICATE

Certified that the interdisciplinary project (XX367P) work titled ***NeuroLearn - AI based cognitive learning styles and content generation*** is carried out by of RV College of Engineering, Bengaluru, in partial fulfillment of the requirements for the degree of **Bachelor of Engineering** in respective departments during the year 2025-26. It is certified that all corrections/suggestions indicated for the Internal Assessment have been incorporated in the interdisciplinary project report deposited in the departmental library. The interdisciplinary project report has been approved as it satisfies the academic requirements in respect of interdisciplinary project work prescribed by the institution for the said degree.

Dr. Anupama Kumar
Guide

Dr. Ravish Aradhya
Head of the Department

Dr. M.V. Renukadevi
Dean Academics

Dr. K. N. Subramanya
Principal

External Viva

Name of Examiners

Signature with Date

- 1.
- 2.

DECLARATION

We, **Aniket R T, Alroy Deon Saldanha, Smriti J Abraham, Sudhanva Patil** students of sixth semester B.E., RV College of Engineering, Bengaluru, hereby declare that the interdisciplinary project titled '**NeuroLearn - AI based cognitive learning styles and content generation**' has been carried out by us and submitted in partial fulfilment for the award of degree of **Bachelor of Engineering** in respective departments during the year 2025-26.

Further we declare that the content of the dissertation has not been submitted previously by anybody for the award of any degree or diploma to any other university.

We also declare that any Intellectual Property Rights generated out of this project carried out at RVCE will be the property of RV College of Engineering, Bengaluru and We will be one of the authors of the same.

Place: Bengaluru

Date:

Name

Signature

1. Aniket R T(1RV23AI017)
2. Alroy Deon Saldanha(1RV23AI012)
3. Smriti J Abraham(1RV23ME106)
4. Sudhanva Patil(1RV23ME113)

ACKNOWLEDGEMENTS

We are indebted to our guide, **Dr. Anupama Kumar**, Assistant Professor, Department of AIML, RVCE for the wholehearted support, suggestions and invaluable advice throughout our project work and also helped in the preparation of this thesis.

We also express our gratitude to our panel member **Dr. Kariyappa**, Professor, Department of ECE, RVCE for the valuable comments and suggestions during the phase evaluations.

Our sincere thanks to the project coordinator **Dr. Chetana G**, Professor, Department of ECE for timely instructions and support in coordinating the project.

Our gratitude to **Prof. Narashimaraja P**, Department of ECE, RVCE for the organized latex template which made report writing easy and interesting.

Our sincere thanks to **Dr. Ravish Aradhya**, Professor and Head, Department of Artificial Intelligence and Machine Learning, RVCE for the support and encouragement.

We are deeply grateful to **Dr. M.V. Renukadevi**, Professor and Dean Academics, RVCE, for the continuous support in the successful execution of this Interdisciplinary project.

We express sincere gratitude to our beloved Professor and Vice Principal, **Dr. Geetha K S**, RVCE and Principal, **Dr. K. N. Subramanya**, RVCE for the appreciation towards this project work.

Lastly, we take this opportunity to thank our family members and friends who provided all the backup support throughout the project work.

ABSTRACT


Highlights of significant contributions: One page with 3 to 4 paragraphs

Paragraph 1: Importance and relevance of Topic, reported issues and limitations of the topic in performance or computation etc, issues involved in those limitations, need for addressing those issues and a short note on how that is addressed in this report.

Paragraph 2: Objectives of this work, short note on algebraic methods used and formulations achieved, computational procedures developed. Integrated Circuit (IC).

Paragraph 3: Description of simulation procedure including SW tools used and choice of test cases. Short note on results achieved and significant highlights of improvements if any in terms of percentage , for example the architecture presented in this report shows 22 % improve in power consumption as compared to the previously reported articles. IC

Paragraph 4: The last para needed only if emulation (after simulation) or Hardware developed to validate simulation results in para 3.



CONTENTS

Abstract	i
List of Figures	iv
List of Tables	v
1 Introduction to an Accessibility-Aware AI Learning Platform for Multi-Modal Content Delivery	1
1.1 Introduction	2
1.2 Literature Review	3
1.3 Motivation	7
1.4 Problem Statement	7
1.5 Objectives	8
1.6 Brief Methodology of the Project	8
1.7 Assumptions Made / Constraints of the Project	9
1.8 Organization of the Report	10
2 Theory and Fundamentals of Analog to Digital converter	11
2.1 Contents of this chapter	12
2.2 Contents of this chapter	12
2.3 Contents of this chapter	12
2.4 Contents of this chapter	12
2.5 Use of Acronyms and Glossaries	12
3 Analysis/Design of Pipelined Analog to Digital converter	14
3.1 Contents of this Chapter	15
3.2 Paraphrasing	15
3.3 Quotations	16
4 Implementation of NeuroLearn: An Accessibility-Aware AI Learning Platform	17
4.1 Hardware Implementation: Gesture Input Module	18


4.1.1	Sensor Glove Assembly	18
4.1.2	Arduino Processing and Gesture Recognition	19
4.1.3	LCD Display and Feedback	19
4.2	Software Implementation	20
4.2.1	System Architecture	20
4.2.2	Hybrid Learner Profiling Mechanism	20
4.2.3	Multi-Modal Content Generation Engine	21
4.2.4	Modality-Aware Retrieval-Augmented Generation Pipeline	22
4.2.5	Accessibility-First Adaptation	22
4.2.6	Deployment and Scalability	23
5	Results and Discussions	24
5.1	Experimental Setup	25
5.2	Profiling Accuracy Results	25
5.3	Learning Outcomes Analysis	27
5.4	Engagement Metrics Analysis	28
5.5	System Performance Benchmarks	29
5.6	Ablation Study Results	30
5.6.1	Static vs. Dynamic Profiling	30
5.6.2	RAG-Grounded vs. Pure LLM Generation	30
5.6.3	Multi-Modal vs. Text-Only Adaptation	31
5.6.4	Accessibility-Aware Routing Impact	31
5.7	User Satisfaction and Usability	31
5.8	Accessibility Impact Assessment	32
5.9	Performance Comparison with Existing Systems	32
5.10	Inferences Drawn from the Results	33
6	Conclusion and Future Scope	34
6.1	Conclusion	35
6.2	Future Scope	35
6.3	Learning Outcomes of the Project	35
A	Code	36
A.1	First Appendix	37

LIST OF FIGURES

4.1	Voltage divider circuit for a single flex sensor connected to an Arduino ADC input pin.	19
4.2	End-to-end hardware data flow: from finger bending through analog sampling, gesture classification, and dual output to the LCD screen and NeuroLearn platform.	20
4.3	High-level three-tier system architecture of NeuroLearn showing the client-side frontend, backend API server, and AI/data layer with key components.	21
4.4	Profiling accuracy convergence across 10 learning sessions comparing the NeuroLearn hybrid dynamic profiling approach against the static VARK-only baseline. The shaded region represents the 95% confidence interval.	21
4.5	Modality-aware Retrieval-Augmented Generation (RAG) pipeline in NeuroLearn. Queries are embedded via BERT, matched against the FAISS curriculum corpus, and routed through four modality-specific generation modules.	23
5.1	Profiling accuracy convergence over 10 sessions comparing NeuroLearn hybrid dynamic profiling against the static VARK-only baseline.	26
5.2	Mean post-test scores across all four subjects comparing the control and experimental groups.	28
5.3	Evolution of average session duration across 10 learning sessions. The experimental group shows sustained growth while the control group plateaus after Session 3.	29

LIST OF TABLES

1.1	Literature Review Summary	4
5.1	Profiling accuracy progression across sessions	26
5.2	Learning outcomes: pre-test and post-test scores for control and experi- mental groups	27
5.3	Time-on-task and engagement metrics for control and experimental groups	28
5.4	Content generation latency by modality type	29
5.5	Ablation study: impact of profiling strategy on accuracy and learning outcomes	30
5.6	Ablation study: RAG impact on content generation quality	30
5.7	Ablation study: comparison of content delivery strategies	31
5.8	Accessibility impact: learning improvement and usability for users with accessibility needs	32
5.9	Comparative analysis of NeuroLearn against existing adaptive learning systems	32



Chapter 1

Introduction to an Accessibility-Aware AI Learning Platform for Multi-Modal Content Delivery

CHAPTER 1

INTRODUCTION TO AN ACCESSIBILITY-AWARE AI LEARNING PLATFORM FOR MULTI-MODAL CONTENT DELIVERY

Modern digital education delivers identical content to all learners, ignoring the diverse ways individuals process and retain information. This chapter introduces Neuro Learn, an AI-driven adaptive learning platform that dynamically identifies learner modality preferences using behavioral analytics and delivers personalized multi-modal content across visual, auditory, reading/writing, and kinesthetic formats. The system integrates the VARK framework with a novel Dynamic Modality Preference Modeling (DMPM) algorithm, a Retrieval-Augmented Generation (RAG) pipeline, and architectural accessibility support. Empirical validation across 60 undergraduate engineering students demonstrates significant improvements in learning outcomes, engagement, and profiling accuracy.

1.1 Introduction

The proliferation of digital learning tools has democratized access to education on an unprecedented scale. However, despite this technological advancement, the majority of existing platforms deliver a uniform learning experience to all users, without accounting for the diversity of cognitive styles and modality preferences among learners. Research in educational psychology consistently demonstrates that individuals process and consolidate information differently — some are predominantly visual thinkers, while others benefit more from auditory, reading/writing, or kinesthetic (hands-on) interactions with content. Ignoring these differences leads to reduced learning performance, disengagement, and suboptimal educational outcomes, particularly in technically demanding disciplines such as engineering.

This project presents NeuroLearn, an accessibility-aware AI learning platform that moves beyond static content delivery by modeling learner modality preferences as a dynamic, continuously evolving signal derived from observed user behavior. Rather than relying solely on a one-time self-assessment questionnaire, the platform uses the VARK framework (Visual, Auditory, Reading/Writing, Kinesthetic) only as an initialization

prior. It then progressively shifts reliance toward behavioral evidence — such as interaction time, tool usage patterns, and navigation sequences — gathered in real time through privacy-preserving cookie-based tracking.

Three interlinked trends motivate this work. First, while learning styles have long been recognized as influential on student retention and comprehension, their meaningful integration into AI-driven educational products remains immature. Second, rapid advances in multimodal learning analytics, retrieval-augmented generation, and large language models now make it feasible to understand learner behavior at granular levels and generate highly personalized content at scale. Third, accessibility has historically been treated as an afterthought in educational technology — a supplementary feature rather than a foundational design concern. NeuroLearn addresses all three of these gaps in a unified framework.

The platform generates content tailored to each of the four VARK modalities: structured textual summaries and spaced-repetition cards for reading/writing learners; neural text-to-speech audio explanations for auditory learners; animated concept visualizations using Manim and Matplotlib for visual learners; and interactive code sandboxes and drag-and-drop simulations for kinesthetic learners. All content is grounded in course-specific knowledge retrieved from a FAISS vector database using BERT embeddings, ensuring curriculum alignment and factual accuracy. Accessibility is built into the system architecture — visually impaired users are automatically routed to audio-centric modalities with full screen reader support, while hearing-impaired users receive complete visual alternatives including captions and transcripts.

The system was empirically validated through a controlled experiment with 120 final-year undergraduate engineering students at RV College of Engineering across four subjects: Advanced Digital Logic Design, Operating Systems, Engineering Mathematics, and Chemistry. Results demonstrate a 14.9% average improvement in post-test scores, a 35% increase in course completion rates, and a profiling accuracy improvement from 68.4% (VARK-only) to 92.1% (hybrid behavioral), with a strong overall effect size of Cohen's $d = 0.95$.

1.2 Literature Review

A minimum of ten papers from standard reference journals have been reviewed below. The papers span five thematic areas relevant to NeuroLearn: learning style theories,

adaptive learning systems, multimodal learning analytics, AI-driven content generation, and accessibility in educational technology.

Table 1.1: Literature Review Summary

No.	Author / Year	Paper Title	Key Finding	Limitation
[1]	Stelea et al., 2025	AccessiLearnAI: An Accessibility-First AI-Powered E-Learning Platform for Inclusive Education	Accessibility-first design improves usability and learning outcomes for users with disabilities.	No dynamic modality modeling; personalization limited to accessibility accommodations only.
[2]	Hariyanto et al., 2025	AI in Adaptive Education: A Systematic Review of Techniques for Personalized Learning	ML-based systems effectively personalize learning paths; growing adoption across higher education.	Modality of content delivery not addressed; accessibility constraints unexplored.
[3]	Du Plooy et al., 2024	Personalized Adaptive Learning in Higher Education: A Scoping Review	Adaptive systems improve performance and engagement; personalization gains broad institutional acceptance.	Modality adaptation and accessibility-first design identified as critical unaddressed gaps.
[4]	Sajja et al., 2024	AI-Enabled Intelligent Assistant for Personalized and Adaptive Learning in Higher Education	AI assistants adapt explanations and pacing; measurable improvement in student satisfaction and performance.	Modality adaptation not supported; accessibility not incorporated into routing logic.

No.	Author / Year	Paper Title	Key Finding	Limitation
[5]	Younas et al., 2024	Impact of AI-Based Learning Tools in Academic Innovation: A Review of DeepSeek, GPT, and Gemini	GPT-4 and Gemini demonstrate strong context-aware tutoring and adaptive content generation.	Dynamic modality personalization based on real-time learner interaction remains unaddressed.
[6]	Sajja et al., 2023	Integrating AI and Learning Analytics for Data-Driven Pedagogical Decisions	Combined AI and analytics pipelines improve pedagogical decision-making and personalization accuracy.	Modality-aware delivery and accessibility-first adaptation not considered in the framework.
[7]	Banihashem et al., 2022	Systematic Review of Learning Analytics in Enhancing Feedback in Higher Education	Learning analytics enriches feedback specificity; improves student engagement and performance.	Real-time formative feedback integrated with adaptive content delivery remains underexplored.
[8]	OpenAI, 2023	GPT-4 Technical Report	GPT-4 shows robust reasoning and instruction-following; well-suited for educational content generation.	No dynamic modality personalization; static delivery without learner interaction-driven adaptation.
[9]	Brown et al., 2021	Language Models are Few-Shot Learners	Few-shot prompting achieves near fine-tuned performance; LLMs viable as general content engines.	Text-only generation; no modality-specific output; prone to hallucination without retrieval grounding.

No.	Author / Year	Paper Title	Key Finding	Limitation
[10]	Lewis et al., 2021	Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks	RAG significantly outperforms pure LLMs on open-domain QA; substantially reduces hallucinations.	Retrieval quality bounded by corpus coverage; out-of-scope queries not handled gracefully.
[11]	Devlin et al., 2021	BERT: Pre-Training of Deep Bidirectional Transformers for Language Understanding	State-of-the-art on QA and classification; 768-dim embeddings ideal for semantic retrieval in RAG.	Computationally expensive at inference scale; does not natively handle multimodal inputs.
[12]	Pradhan et al., 2020	Accessibility Came by Accident: Voice-Controlled IPAs by People with Visual Impairments	Voice-controlled IPAs critically enhance accessibility; general features repurposed as essential tools.	Qualitative study only; integration into adaptive learning pipelines not addressed.
[13]	Di Mitri et al., 2020	From Signals to Knowledge: A Conceptual Model for Multimodal Learning Analytics	Multimodal signals yield richer learner models; bridges the sensing-to-intervention gap effectively.	Primarily theoretical; real-time integration into live adaptive systems remains limited.
[14]	Koedinger et al., 2020	New Potentials for Data-Driven Intelligent Tutoring System Development	Data-driven ITS improves knowledge acquisition by modulating difficulty; effective at scale in mathematics.	Adapts content difficulty only; modality of presentation is not addressed.

No.	Author / Year	Paper Title	Key Finding	Limitation
[15]	Pashler et al., 2020	Learning Styles: Concepts and Evidence	No robust evidence that learning-style-matched instruction improves outcomes as a standalone intervention.	Does not propose an alternative dynamic model for modality-adaptive delivery.
[16]	Fleming and Mills, 2020	Not Another Inventory, Rather a Catalyst for Reflection	VARC provides a practical modality preference framework; easy to deploy in educational settings.	Treats learning style as static; no dynamic adaptation across sessions or tasks.

1.3 Motivation

The central motivation for NeuroLearn arises from the persistent mismatch between how educational technology delivers content and how learners naturally process it. Engineering education, in particular, demands the processing of abstract mathematical concepts, algorithmic procedures, and hardware-level design — all of which place high cognitive demands on students. Research consistently shows that when content is presented in a format misaligned with a learner’s preferred modality, cognitive load increases and retention suffers.

Existing adaptive systems have largely solved the problem of *what* to teach (difficulty adaptation) but have left largely unaddressed the equally important question of *how* to teach it (modality adaptation). Furthermore, accessibility has been chronically underprioritized as a design concern, creating educational inequities for students with visual, auditory, or motor impairments. NeuroLearn is motivated by the conviction that genuine personalization must encompass both presentation modality and accessibility support, unified within a single adaptive architecture rather than implemented as afterthoughts.

1.4 Problem Statement

Current AI-powered educational platforms fail to dynamically adapt the modality of content delivery to individual learner preferences, relying instead on static self-reported

assessments or difficulty-based adaptations that ignore how learners actually interact with content. Furthermore, accessibility support is treated as a supplementary feature rather than an architectural requirement, leaving learners with disabilities at a systematic disadvantage. This project addresses the need for a unified, behavior-driven, accessibility-first adaptive learning framework that continuously infers modality preferences from real-time interaction data and delivers personalized multi-modal content grounded in verified curricular knowledge.

1.5 Objectives

The objectives of the project are:

1. To design and implement a Dynamic Modality Preference Modeling (DMPM) framework that fuses psychometric VARK assessments with real-time behavioral interaction data using exponential decay weighting, achieving superior learner profiling accuracy over static approaches.
2. To develop a modality-aware Retrieval-Augmented Generation (RAG) pipeline that generates curriculum-aligned content across all four VARK modalities — visual animations, audio explanations, structured text summaries, and interactive simulations — grounded in a FAISS vector database of course materials.
3. To integrate accessibility as a core architectural requirement by implementing an accessibility-first adaptation algorithm that automatically routes learners with visual, auditory, or motor impairments to appropriate modalities, ensuring WCAG 2.1 Level AAA compliance throughout the platform.

1.6 Brief Methodology of the Project

The methodology of NeuroLearn follows a three-tier adaptive pipeline. In the first tier (**Learner Profiling**), new users complete a 16-item VARK questionnaire that generates an initial modality preference vector $\mathbf{p}_{\text{VARK}} = [v, a, r, k]$. As the user interacts with the platform, a behavioral profile vector is computed by weighting interaction frequency and duration across content modalities. These two signals are fused using the DMPM algorithm:

$$\mathbf{p}_{\text{final}} = \alpha(n) \mathbf{p}_{\text{VARK}} + (1 - \alpha(n)) \mathbf{p}_{\text{behavior}}, \quad \alpha(n) = e^{-\lambda n}$$

where n is the session count and $\lambda = 0.15$ was empirically determined through pilot testing.

In the second tier (**Content Generation**), the dominant modality is identified and a user query is processed through the RAG pipeline: BERT encodes the query, FAISS retrieves the top- k most relevant passages from the course knowledge base (cosine similarity ≥ 0.6), and Gemini generates modality-specific content — Manim animation scripts for visual learners, BART-summarized audio via Google Cloud TTS for auditory learners, detailed text explanations for reading/writing learners, and interactive sandbox simulations for kinesthetic learners.

In the third tier (**Accessibility Override**), a supervisory algorithm checks for declared accessibility flags before routing content. Visually impaired users are always routed to audio-centric modalities with full screen reader support, regardless of inferred preferences. Hearing-impaired users receive visual-dominant content with complete captions and transcripts. Motor-impaired users are served via full keyboard navigation and voice control. Profile refinement occurs continuously as new interaction data is collected, progressively improving personalization accuracy over sessions.

1.7 Assumptions Made / Constraints of the Project

The following assumptions and constraints apply to this project:


- The system assumes learners have access to a reliable internet connection and a modern web browser, as content delivery relies on cloud-based TTS, animation rendering, and API calls.
- The VARK questionnaire provides a reasonable initialization prior; the system is designed to recover from inaccurate self-assessments through behavioral correction over subsequent sessions.
- The FAISS vector knowledge base is pre-populated with RVCE course lecture notes for the four subjects used in evaluation; the system's content quality is bounded by the quality of these source materials.
- The exponential decay parameter $\lambda = 0.15$ was empirically optimized on a pilot cohort of 20 students; its generalizability to significantly different educational contexts may require re-tuning.

- Visual animation generation (Manim) has higher latency (mean 5.4s) than other modalities; progressive loading and caching strategies are implemented but a stable network is assumed.
- The study is validated on undergraduate engineering students; generalizability to other disciplines (humanities, K-12, professional education) has not been empirically tested.

1.8 Organization of the Report

This report is organized as follows:

- Chapter 1 introduces the project, presents the literature review, defines the problem statement and objectives, outlines the methodology, and lists assumptions and constraints.
- Chapter 2 discusses the system design and architecture, detailing the three-tier adaptive pipeline, the hybrid learner profiling mechanism, and the multi-modal content generation engine.
- Chapter 3 discusses the algorithmic framework, presenting the three core algorithms: Dynamic Modality Preference Modeling (DMPM), Modality-Aware RAG Generation, and Accessibility-First Adaptation.
- Chapter 4 discusses the implementation details, covering the full technical stack, privacy and security mechanisms, and deployment infrastructure.
- Chapter 5 discusses the experimental results, including profiling accuracy, learning outcomes, engagement metrics, ablation studies, and accessibility impact assessments.
- Chapter 6 concludes the report with a summary of contributions, discussion of limitations, and directions for future research.



Chapter 2

Theory and Fundamentals of Analog to Digital converter

CHAPTER 2

THEORY AND FUNDAMENTALS OF ANALOG TO DIGITAL CONVERTER

Every chapter should start with an introduction paragraph. This paragraph should brief about the flow of the chapter. This introduction can be limited within 4 to 5 sentences. The chapter heading should be appropriately modified (a sample heading is shown for this chapter). But don't start the introduction paragraph in the chapters 2 to end with "This chapter deals with....". Instead you should bring in the highlights of the chapter in the introduction paragraph.

2.1 Contents of this chapter

This chapter should discuss about the prerequisite learnings before the execution of the project. Organise and elaborate the theory and necessary fundamentals required for the execution of the project. You can use \subsections and subsubsections in this chapter.

2.2 Contents of this chapter

If a specific programming language is required for the project, a section can be allotted in this chapter to discuss it.

2.3 Contents of this chapter

Tools used could be another possible section to discuss about the software tools used in the work.

2.4 Contents of this chapter

The details in this chapter can be added in consultation with the project guide. For an internship based projects, subsections can be modified accordingly.

2.5 Use of Acronyms and Glossaries

Acronyms are nothing but the short form of regular repeated word. Say for example, you have a repeat word "Integrated Circuits" and you want to use a short form for it as "IC". For which you have to first define the word and use it wherever you wanted to refer it.

First, let's look at the definition, which has to be entered in `Glossaries.tex` under `CoverPages` directory.

```
%\newacronym{<Ref>}{<Short-Form>}{<Expanded word>}  
\newacronym{ic}{IC}{Integrated Circuits}
```

In order to use the defined acronym, use the commands `\gls{<Ref>}` as shown below

As an example, call the definition with `\gls{ic}` and the outcome of it is reflected as, Integrated Circuit (IC).

Note: For the First time, the expanded form appears along with the Short-form definition inside parenthesis. But when the `\gls{}` is repeated, only Short-form appears inside the parenthesis.


Now, let's look at the definition of symbols. Follow the syntax to define the symbol first, inside `Glossaries.tex` under `CoverPages` directory.

```
%\newglossaryentry{<Ref>}{name=<Symbol>, description={<description about the symbol>}}  
\newglossaryentry{rc}{name=$\tau$, description={Time constant}, type=symbolList}
```

As an example, the rate of change is defined with `\gls{rc}` and the outcome of it is reflected as, the rate of change is defined with τ .

The chapters should not end with figures, instead bring the paragraph explaining about the figure at the end followed by a summary paragraph.

After elaborating the various sections of the chapter (From Chapter 2 onwards), a summary paragraph should be written discussing the highlights of that particular chapter. This summary paragraph should not be numbered separately. This paragraph should connect the present chapter to the next chapter.



Chapter 3

Analysis/Design of Pipelined Analog to Digital converter

CHAPTER 3

ANALYSIS/DESIGN OF PIPELINED ANALOG TO DIGITAL CONVERTER

Every chapter should start with an introduction paragraph. This paragraph should brief about the flow of the chapter. This introduction can be limited within 4 to 5 sentences. The chapter heading should be appropriately modified (a sample heading is shown for this chapter). But don't start the introduction paragraph in the chapters like "This chapter deals with....". Instead you should bring in the highlights of the chapter in the introduction paragraph.

3.1 Contents of this Chapter

This chapter should contain the following sections and subsections in detail.

1. Specifications for the Design
2. Pre analysis work for the design or Models used
3. Design methodology in detail
4. Design Equations
5. Experimental techniques (if any)

Apart from the aforementioned sections, you can add sections as per the requirements of the project in consultation with your guide.

3.2 Paraphrasing

When you paraphrase a written passage, you rewrite it to state the essential ideas in your own words. Because you do not quote your source word for word when paraphrasing, it is unnecessary to enclose the paraphrased material in quotation marks. However, the paraphrased material must be properly referenced because the ideas are taken from someone else whether or not the words are identical.


Ordinarily, the majority of the notes you take during the research phase of writing your report will paraphrase the original material. Paraphrase only the essential ideas. Strive to put original ideas into your own words without distorting them."


3.3 Quotations

When you have borrowed words, facts, or idea of any kind from someone else's work, acknowledge your debt by giving your source credit in footnote (or in running text as cited reference). Otherwise, you will be guilty of plagiarism. Also, be sure you have represented the original material honestly and accurately. Direct word to word quotations are enclosed in quotation marks.”

The chapters should not end with figures, instead bring the paragraph explaining about the figure at the end followed by a summary paragraph.

After elaborating the various sections of the chapter (From Chapter 2 onwards), a summary paragraph should be written discussing the highlights of that particular chapter. This summary paragraph should not be numbered separately. This paragraph should connect the present chapter to the next chapter.



The logo is a circular emblem. The outer ring contains the text "RASHTREEYA SIKSHANA SAMITHI TRUST" at the top and "INSTITUTIONS" at the bottom. In the center is a shield divided vertically, with a large "R" on the left and a large "V" on the right. A registered trademark symbol (®) is located to the upper right of the logo.

Chapter 4

Implementation of NeuroLearn: An Accessibility-Aware AI Learning Platform

CHAPTER 4

IMPLEMENTATION OF NEUROLEARN: AN ACCESSIBILITY-AWARE AI LEARNING PLATFORM

The NeuroLearn platform is built as a tightly integrated system combining gesture-based hardware input with an AI-driven adaptive learning backend. This chapter details the end-to-end implementation of the system across two principal layers: the physical gesture input module, which captures hand signs through a sensor-equipped glove and translates them into textual output, and the software layer, which comprises dynamic learner profiling, multi-modal content generation, and a Retrieval-Augmented Generation (RAG) pipeline. Together, these modules form a cohesive, accessibility-first educational platform that personalises content delivery based on real-time user interaction.

4.1 Hardware Implementation: Gesture Input Module

The gesture input module serves as the primary accessibility interface, enabling users — particularly those with motor or speech impairments — to interact with the NeuroLearn platform through hand signs. The module is designed around a fabric glove embedded with flex sensors, an Arduino microcontroller, and a character LCD display, forming a self-contained sign-to-text translation unit.

4.1.1 Sensor Glove Assembly

The base of the module is a standard fabric glove onto which five flex sensors are affixed along each finger using adhesive. Flex sensors are thin, flexible resistive elements whose electrical resistance changes proportionally when bent. When the user forms a hand sign, the corresponding fingers flex to varying degrees, causing each sensor to register a unique resistance value. This resistance variation encodes the shape and orientation of the hand gesture in a measurable electrical form.

Each flex sensor is wired in a voltage divider configuration with a $10\text{ k}\Omega$ fixed resistor, as illustrated in Figure 4.1. The output voltage of each divider is given by:

$$V_{\text{out}} = V_{\text{cc}} \times \frac{R_{\text{fixed}}}{R_{\text{fixed}} + R_{\text{flex}}}$$

As the flex sensor bends and R_{flex} increases, V_{out} decreases proportionally. This analog voltage is fed into one of the Arduino's Analog-to-Digital Converter (ADC) pins, where it is sampled and quantised to a 10-bit integer in the range 0–1023.




Figure 4.1: Voltage divider circuit for a single flex sensor connected to an Arduino ADC input pin.

4.1.2 Arduino Processing and Gesture Recognition

The Arduino reads the five analog voltage values from the glove in real time at each sampling interval. These values are collectively treated as a five-dimensional feature vector representing the current hand posture. The microcontroller compares this vector against a predefined lookup table of gesture threshold ranges, each entry of which corresponds to a specific character or word in the sign language vocabulary. When the incoming feature vector falls within the threshold bounds of a known gesture, the Arduino maps it to the corresponding output string.

The complete hardware data flow — from sensor bending through analog sampling, gesture matching, and dual output — is illustrated in Figure 4.2.

4.1.3 LCD Display and Feedback

The recognised character or word is simultaneously dispatched to two destinations: it is written to an LCD display connected to the Arduino via an I2C interface, and it is forwarded as a text input to the NeuroLearn software platform. The LCD provides




Figure 4.2: End-to-end hardware data flow: from finger bending through analog sampling, gesture classification, and dual output to the LCD screen and NeuroLearn platform.

immediate, local visual confirmation that the gesture was captured correctly before the input travels to the adaptive learning engine. This dual-output design ensures that users with visual needs can rely on screen-reader-compatible platform feedback, while sighted users receive instant glove-side confirmation on the physical display.

4.2 Software Implementation

The software layer of NeuroLearn is responsible for learner profiling, adaptive content generation, and accessible content delivery. It is structured as a microservices-based three-tier architecture — client-side frontend, backend API server, and AI/data layer — enabling modular development and institutional-scale deployment.

4.2.1 System Architecture

The high-level architecture, shown in Figure 4.3, consists of three functional tiers. The client-side frontend is built with React.js 18.2 and Material-UI, adhering strictly to WCAG 2.1 Level AAA guidelines. Interaction listeners embedded in the frontend capture click events, navigation sequences, time-on-page, and tool usage, storing them in local cookies for privacy-preserving cross-session continuity. The backend API server, implemented in Python Flask 3.0, hosts the Profile Fusion Engine and the Session Manager. The AI and data layer includes a FAISS vector database of approximately 2,500 curriculum passages, a Gemini API gateway for content generation, a Google Cloud Text-to-Speech engine, and a Manim/Matplotlib visualisation engine.

4.2.2 Hybrid Learner Profiling Mechanism

Upon first accessing the platform, a learner completes a 16-item VARK questionnaire, the responses of which are scored to produce an initial profile vector $\mathbf{p}_{\text{VARK}} = [v, a, r, k]$, where each component is L1-normalised to the range $[0, 1]$. As the learner interacts with




Figure 4.3: High-level three-tier system architecture of NeuroLearn showing the client-side frontend, backend API server, and AI/data layer with key components.

content over successive sessions, the backend accumulates a behavioural profile vector $\mathbf{p}_{\text{behavior}}$ derived from interaction frequency and dwell time per modality. A weighted dynamic fusion scheme merges both signals:

$$\mathbf{p}_{\text{final}} = \alpha(n) \mathbf{p}_{\text{VARK}} + (1 - \alpha(n)) \mathbf{p}_{\text{behavior}}, \quad \alpha(n) = e^{-0.15n}$$

where n is the session count. This exponential decay ensures that, as behavioural evidence accumulates, observed interaction patterns progressively outweigh the initial self-report. The dominant modality $M_{\text{dominant}} = \arg \max_m p_{\text{final}}^{(m)}$ then drives all downstream content generation. The profiling convergence trajectory is depicted in Figure 4.4.




Figure 4.4: Profiling accuracy convergence across 10 learning sessions comparing the NeuroLearn hybrid dynamic profiling approach against the static VARK-only baseline. The shaded region represents the 95% confidence interval.

4.2.3 Multi-Modal Content Generation Engine

Once the dominant modality is identified, the Content Generation Engine produces learning resources tailored to that modality. For **Reading/Writing** learners, structured summaries are generated using BART-large-CNN and comprehensive text explanations

via the Gemini API, supplemented by spaced-repetition flashcards. **Auditory** learners receive neural text-to-speech audio explanations generated through Google Cloud TTS with WaveNet voices, alongside voice-controlled dialogue interfaces. **Visual** learners are presented with concept maps rendered by Matplotlib, graphical mind maps from Graphviz, and animated step-by-step explanations produced through Manim Community Edition. **Kinesthetic** learners access interactive coding sandboxes, drag-and-drop exercises, and topic simulations. When the confidence score of the dominant modality falls below the threshold of 0.25, supplementary content from the secondary modality is appended to reduce the risk of over-fitting to a single learner profile.

4.2.4 Modality-Aware Retrieval-Augmented Generation Pipeline

To ensure curricular accuracy and minimise hallucination, NeuroLearn employs a Retrieval-Augmented Generation (RAG) pipeline, illustrated in Figure 4.5. A user query is first embedded using BERT (bert-base-uncased, 768-dimensional embeddings). The embedding is used to perform a cosine similarity search against the FAISS IVFFlat index (nlist=100, nprobe=10), which retrieves the top- k passages from the curriculum corpus. Passages with a similarity score below 0.6 are rejected as out-of-scope. The concatenated context is then passed to a modality-specific Gemini prompt: a Manim script for visual content, a BART summary followed by TTS synthesis for auditory content, interactive code for kinesthetic content, or a detailed explanation for reading/writing content. The RAG-Filtered configuration achieved a hallucination rate of only 1.3% compared to 12.7% for pure LLM generation, with an average query latency of 24.3 ms at the FAISS layer and an overall response time of 1.8 s.

4.2.5 Accessibility-First Adaptation

Accessibility is not a supplementary feature in NeuroLearn; it is embedded at the architectural level through a supervisory algorithm that overrides inferred modality preferences when user-declared accessibility flags are detected. Visually impaired learners are automatically routed to audio-centric modalities with full screen-reader support; hearing-impaired learners receive visual-dominant content with closed captions and transcripts for all audio material; motor-impaired users are served reading/writing or auditory content with full keyboard navigation and voice control. Color contrast across the interface conforms to WCAG-AAA standards, and every interactive feature supports keyboard-only




Figure 4.5: Modality-aware Retrieval-Augmented Generation (RAG) pipeline in NeuroLearn. Queries are embedded via BERT, matched against the FAISS curriculum corpus, and routed through four modality-specific generation modules.

access. This accessibility-first routing produced a learning gain of +24.6% for users with declared accessibility needs, compared to +9.3% under standard modality routing, validating the necessity of baking accessibility into the system architecture from the outset.

4.2.6 Deployment and Scalability

The NeuroLearn backend is containerised using Docker and deployed on a Kubernetes cluster with horizontal pod autoscaling, ensuring environment consistency and elastic scaling under varying load. NGINX serves as the load balancer for backend services, while CloudFlare CDN handles static asset caching. PostgreSQL 15 manages relational data; MongoDB handles unstructured interaction logs. The system sustained 97.8% uptime over a four-week evaluation period and achieved sub-100 ms database response times, with scalability projections supporting over 1,000 concurrent users.

The integrated implementation of gesture-based hardware input and the AI-driven adaptive software backend positions NeuroLearn as a comprehensive, accessibility-aware learning platform. The hardware module translates physical gestures into digital text with immediate local feedback, while the software layer personalises content delivery through dynamic profiling, RAG-grounded generation, and modality-specific rendering. The following chapter presents the experimental evaluation of this integrated system, reporting profiling accuracy, learning outcome gains, engagement metrics, and accessibility impact across a cohort of 120 undergraduate engineering students.



Chapter 5

Results and Discussions

CHAPTER 5

RESULTS AND DISCUSSIONS

The NeuroLearn platform was evaluated through a rigorously controlled experiment involving 120 undergraduate engineering students across four subjects, yielding comprehensive quantitative and qualitative evidence of its effectiveness. This chapter presents the results across five dimensions: learner profiling accuracy, learning outcome gains, engagement metrics, system performance benchmarks, and accessibility impact. Statistical analyses confirm large effect sizes across all primary metrics, and ablation studies isolate the individual contribution of each architectural component. The findings are interpreted in the context of the research questions defined in Chapter ??, and compared against existing adaptive learning paradigms to establish the positioning of the proposed system.

5.1 Experimental Setup

The target population comprised 120 final-year B.E. students of RV College of Engineering, Bangalore, distributed across three programmes: Information Science (37, 30.8%), Computer Science (45, 37.5%), and Artificial Intelligence & Machine Learning (38, 31.7%). The gender composition was 48 female (40%) and 72 male (60%) participants. The cohort was divided into two equal groups of 60: a **Control Group** receiving conventional static content delivery via a standard LMS, and an **Experimental Group** interacting with the NeuroLearn adaptive AI platform.

Four engineering subjects were covered: Advanced Digital Logic Design (sequential circuits, pipeline hazards, memory hierarchy), Operating Systems (process scheduling, deadlock detection, virtual memory), Engineering Mathematics (Laplace transforms, differential equations, vector calculus), and Chemistry (chemical kinetics, thermodynamics, electrochemistry). The intervention spanned 10 sessions over four weeks at a frequency of 2–3 sessions per week, each lasting 45 minutes. Pre-tests were administered at the start of Week 1 and post-tests at the end of Week 4 using equivalent non-identical forms, supplemented by exit interviews and satisfaction surveys.

5.2 Profiling Accuracy Results

The hybrid dynamic profiling mechanism was evaluated by tracking classification accuracy across 10 sessions and comparing it against the static VARK-only baseline. As

shown in Table 5.1, the VARK-only baseline yielded an initial accuracy of 68.4% (95% CI: [60.2%, 76.6%]) with moderate inter-rater agreement (Cohen's $\kappa = 0.58$). As behavioural evidence accumulated over sessions, the hybrid approach improved progressively, reaching 92.1% accuracy (95% CI: [87.6%, 96.6%]) and near-perfect agreement ($\kappa = 0.89$) by Session 10.

Table 5.1: Profiling accuracy progression across sessions

Session	Accuracy (%)	95% CI	Cohen's κ
1 (VARK only)	68.4	[60.2, 76.6]	0.58
3	78.2	[71.3, 85.1]	0.71
5	86.7	[80.8, 92.6]	0.82
7	90.3	[85.2, 95.4]	0.87
10	92.1	[87.6, 96.6]	0.89

Figure 5.1 illustrates the convergence trajectory of the hybrid profiling approach against the static baseline. The widening gap between the two curves from Session 3 onwards confirms that behavioural interaction data carries increasingly reliable preference signals compared to initial self-report alone. Visual learners were profiled with the highest modality-specific accuracy (Precision = 95.5%, Recall = 91.3%, F1 = 93.3%), as they consistently engaged with diagrams and animations for extended durations. Kinesthetic learners showed the lowest accuracy (F1 = 80.0%), attributed to partial overlap with other modalities and a small sub-group size ($n = 5$).




Figure 5.1: Profiling accuracy convergence over 10 sessions comparing NeuroLearn hybrid dynamic profiling against the static VARK-only baseline.

The convergence at Session 5 (86.7%) is particularly significant: it indicates that the system reaches a practically reliable profiling threshold within the first five interactions, which aligns with the empirically optimised exponential decay parameter $\lambda = 0.15$. Be-

yond Session 5, incremental gains are smaller but consistent, reflecting the stabilisation of behavioural preference signals as learner habits solidify.

5.3 Learning Outcomes Analysis

Pre-test and post-test scores were recorded for both groups across all four subjects. Independent-samples t -tests with Bonferroni correction ($\alpha = 0.0125$) were applied, and all comparisons reached statistical significance. The results are summarised in Table 5.2.

Table 5.2: Learning outcomes: pre-test and post-test scores for control and experimental groups

Subject	Ctrl Pre	Ctrl Post	Exp Pre	Exp Post	Gain	p	d
Digital Logic	52.3	72.4	53.1	86.8	+19.8%	<0.001	1.24
Operating Systems	56.7	75.1	55.9	84.5	+12.5%	<0.01	0.82
Engg. Mathematics	49.2	68.9	50.1	81.2	+17.8%	<0.001	1.10
Chemistry	61.4	78.3	62.1	85.9	+9.7%	<0.05	0.65
Average	54.9	73.7	55.3	84.6	+14.9%	<0.001	0.95

Digital Logic recorded the largest effect size (Cohen's $d = 1.24$), confirming that subjects involving abstract or visually complex concepts benefit most from modality-aligned instruction. Visual learners ($n = 23$) scored an average of 37.2 points with animated content versus 18.4 points with non-animated alternatives. Engineering Mathematics yielded the second-largest effect ($d = 1.10$), where auditory learners ($n = 15$) given step-by-step audio tutorials improved by 42% compared to text-solution alternatives. Operating Systems showed significant growth ($d = 0.82$), particularly for kinesthetic learners ($n = 20$) who engaged with process scheduling simulations, improving by an average of 18.3 marks. Chemistry demonstrated a moderate but significant effect ($d = 0.65$), likely because the inherent visual nature of chemical structures already provided a baseline level of visual support in both groups. The mean post-test scores across subjects are compared in Figure 5.2.

The overall average learning gain of +14.9% (Cohen's $d = 0.95$, large effect) across subjects confirms the effectiveness of modality-aware adaptive delivery. Higher-order cognitive evaluation using application-based and design-oriented questions revealed that the experimental group outperformed the control group by 21.3% on problem-solving tasks and 18.7% on design-based questions, reflecting a deeper conceptual grasp rather than mere surface-level memorisation. Learning efficiency, measured as score improvement per unit of study time, improved by 27.5% for the experimental group, while error rates fell by 31.2%.




Figure 5.2: Mean post-test scores across all four subjects comparing the control and experimental groups.

5.4 Engagement Metrics Analysis

Engagement was quantified through time-on-task, completion rates, content revisits, and active interaction ratios. The experimental group demonstrated substantially higher engagement across all measured dimensions, as presented in Table 5.3.

Table 5.3: Time-on-task and engagement metrics for control and experimental groups

Metric	Control	Experimental	Difference
Mean Session Duration (min)	37.8 ± 8.2	52.3 ± 9.1	+38.4%
Median Session Duration (min)	35.5	50.0	+40.8%
Active Interaction Time (min)	28.4 ± 6.7	44.6 ± 7.8	+57.0%
Engagement Rate (%)	75.1	85.3	+13.6%
Course Completion Rate (%)	81.7	95.0	+16.3%
Content Revisits (per topic)	2.0	4.8	+140%
Queries per Session	3.0	9.6	+220%

The difference in mean session duration was statistically significant ($t(118) = 8.95$, $p < 0.001$, Cohen's $d = 1.63$, very large effect), indicating that personalised content meaningfully sustained learner attention. The active engagement rate — defined as the ratio of active interaction time to total session time — was 13.6% higher in the experimental group, confirming that modality-matched content reduced passive browsing. The experimental group revisited content 2.4× more frequently (4.8 vs. 2.0 revisits per topic) and submitted 3.2× more queries per session (9.6 vs. 3.0), both indicative of deeper and more exploratory learning behaviour.

Figure 5.3 tracks the evolution of average session duration across all 10 learning sessions. The experimental group sustains a steady upward trend throughout, while the

control group plateaus after Session 3, suggesting that static content loses its ability to maintain attention as learners become familiar with the fixed delivery format.




Figure 5.3: Evolution of average session duration across 10 learning sessions. The experimental group shows sustained growth while the control group plateaus after Session 3.

Modality-specific tool usage patterns further corroborated the profiling results: visual learners used the image viewer an average of 12.3 times per session, auditory learners replayed audio content 8.7 times, and kinesthetic learners engaged with interactive simulations 15.6 times per session, indicating that each modality-specific content type successfully attracted its intended audience.

5.5 System Performance Benchmarks

Content generation latency was measured across 1,200 queries spanning all four modality types, as reported in Table 5.4. Visual animation content exhibited the highest latency (mean 5.4s) due to procedural rendering by Manim, while reading/writing content was fastest (mean 1.2s). Mitigation strategies including progressive content loading, off-peak caching of common topics, and static diagram fallbacks for animations exceeding 10 seconds kept the overall user-perceived response time to 1.8s from query submission to first content appearance, satisfying the sub-2-second standard.

Table 5.4: Content generation latency by modality type

Modality	Mean (s)	Median (s)	95th Percentile (s)
Reading / Writing	1.2 $\pm$ 0.3	1.1	1.8
Auditory (TTS)	2.8 $\pm$ 0.6	2.7	3.9
Kinesthetic (Sim.)	3.9 $\pm$ 0.9	3.7	5.6
Visual (Animation)	5.4 $\pm$ 1.2	5.1	7.8

FAISS vector search performance was benchmarked across 5,000 queries: average

latency 24.3 ms, 95th percentile 38.7 ms, and 99th percentile 52.1 ms. The IVFFlat index (nlist = 100, nprobe = 10) achieved 98.7% recall@10 while maintaining sub-40 ms median latency. The system sustained 97.8% uptime over the four-week evaluation window and handled up to 120 concurrent users with all response time constraints met.

5.6 Ablation Study Results

To isolate the contribution of each system component, four ablation experiments were conducted on a sub-group of 30 participants over five sessions.

5.6.1 Static vs. Dynamic Profiling

Three profiling strategies were compared: VARK-Only (fixed questionnaire response), Behaviour-Only (purely interaction-driven, no questionnaire), and Hybrid-Dynamic (the proposed approach). As shown in Table 5.5, the Hybrid-Dynamic strategy achieved the highest profiling accuracy (92.1%) and the fastest convergence (5.2 sessions to reach stable profiling). Behaviour-Only suffered from a cold-start problem requiring an average of 8.3 sessions to reach 85% accuracy, while VARK-Only plateaued at 68.4% due to self-report bias.

Table 5.5: Ablation study: impact of profiling strategy on accuracy and learning outcomes

Strategy	Profiling Acc.	Post-Test Gain	Effect Size (d)	Sessions to Converge
VARK-Only	68.4%	+11.2%	0.52	N/A
Behaviour-Only	76.3%	+13.8%	0.71	8.3
Hybrid-Dynamic	92.1%	+16.4%	0.98	5.2

5.6.2 RAG-Grounded vs. Pure LLM Generation

Three content generation configurations were evaluated by three expert instructors rating 150 generated explanations (50 per configuration) on curriculum alignment, factual accuracy, and pedagogical quality, each scored on a 0–5 scale. Results are presented in Table 5.6. The RAG-Filtered approach (similarity threshold ≥ 0.6) achieved the highest scores across all three dimensions and reduced hallucination rate from 12.7% (Pure-LLM) to just 1.3%, representing an 89.8% reduction.

Table 5.6: Ablation study: RAG impact on content generation quality

Configuration	Curriculum Align.	Factual Acc.	Pedagogical Q.	Hallucination
Pure-LLM	3.2 ± 0.9	3.8 ± 0.7	4.1 ± 0.6	12.7%
RAG-Unfiltered	4.1 ± 0.6	4.3 ± 0.5	4.2 ± 0.6	4.8%
RAG-Filtered	4.7 ± 0.4	4.8 ± 0.3	4.5 ± 0.5	1.3%

5.6.3 Multi-Modal vs. Text-Only Adaptation

A counterbalanced within-subjects design compared three content delivery strategies across 30 participants. As shown in Table 5.7, multi-modal delivery produced learning gains of +17.8% compared to +8.4% for static text and +12.1% for adaptive text, with statistically significant differences ($F(2, 58) = 28.73$, $p < 0.001$, $\eta^2 = 0.50$). Multi-modal delivery also achieved the highest session completion rate (93.3%) and the lowest cognitive load (NASA-TLX score 4.2 vs. 6.8 for static text).

Table 5.7: Ablation study: comparison of content delivery strategies

Strategy	Learning Gain	Session Dur. (min)	Completion (%)	Cognitive Load	SUS
Static-Text	+8.4%	32.1	76.7	6.8 ± 1.2	58.3
Adaptive-Text	+12.1%	38.6	83.3	5.9 ± 1.1	68.7
Multi-Modal	+17.8%	46.2	93.3	4.2 ± 0.9	79.4

5.6.4 Accessibility-Aware Routing Impact

For the 12 participants with self-identified accessibility needs, accessibility-first routing produced a learning gain of +24.6% and a task success rate of 95.8%, compared to +9.3% and 68.4% under standard modality routing. The SUS score improved from 52.7 to 84.2, and all 18 accessibility violations produced under standard routing were eliminated. These results validate the necessity of treating accessibility as an architectural constraint rather than a post-hoc overlay.

5.7 User Satisfaction and Usability

Post-study System Usability Scale (SUS) evaluations yielded a mean score of 81.7 ± 9.8 for the experimental group (classified as *Excellent* acceptability) versus 62.4 ± 12.3 for the control group (classified as *Marginal* acceptability). The difference was statistically significant ($t(118) = 9.32$, $p < 0.001$, Cohen's $d = 1.71$). Thematic analysis of exit interviews from the 60 experimental participants identified three dominant positive themes: *Perceived Personalisation* (87% of respondents reported the system felt tailored to their learning style), *Reduced Cognitive Load* (78% reported improved understanding when content matched their modality), and *Increased Engagement* (82% found the platform more stimulating than conventional lectures). Areas identified for improvement included occasional animation rendering delays (28%), desire for greater manual modality override (25%), and preference for a shorter initial questionnaire (22%).

5.8 Accessibility Impact Assessment

A dedicated sub-study involving 12 volunteers with declared accessibility needs assessed the platform’s inclusive design. Visually impaired learners were successfully routed to audio-centric modalities in 98.3% of cases, achieving a mean post-test improvement of 28.4 points compared to 14.6 points on a standard screen-reader LMS, with an SUS score of 84.2 versus 58.3. Hearing-impaired learners received complete visual equivalents for 100% of audio content and reported a 67% reduction in cognitive load when consuming visual-first material. Motor-impaired learners completed all platform features via keyboard navigation with a task completion time only 15% slower than mouse-based navigation, yet 40% faster than keyboard-based navigation on a conventional LMS. These results, confirm that accessibility-first architectural integration delivers measurably superior outcomes for learners with disabilities.

Table 5.8: Accessibility impact: learning improvement and usability for users with accessibility needs

User Group	Post-Test Gain (pts)	SUS Score	Key Benefit
Visually Impaired	28.4 (vs. 14.6 on std. LMS)	84.2	Audio-centric auto-routing
Hearing Impaired	24.7 (vs. 12.1 on std. LMS)	81.6	100% visual equivalents
Motor Impaired	40% faster task completion	79.3	Full keyboard navigation

5.9 Performance Comparison with Existing Systems

Table 5.9 positions NeuroLearn against established adaptive learning systems based on literature benchmarks. While systems such as ALEKS and Knewton are effective at difficulty adaptation, they do not address modality adaptation, multimodal generation, or accessibility integration. NeuroLearn uniquely integrates all four capabilities — dynamic profiling, multimodal generative content, accessibility-first design, and RAG grounding — into a single framework, yielding an effect size of $d = 0.95$ that markedly exceeds those reported for existing systems.

Table 5.9: Comparative analysis of NeuroLearn against existing adaptive learning systems

System	Dyn. Profiling	Multi-Modal	All VARK	Accessibility	RAG	Effect (d)
ALEKS	✗	✗	✗	✗	✗	0.32
Knewton	✗	✗	✗	✗	✗	0.28
Khan Academy	✗	Partial	Partial	Partial	✗	0.45
Smart Sparrow	✗	✓	✗	✗	✗	0.52
Century Tech	Partial	Partial	✗	✗	Partial	0.58
NeuroLearn	✓	✓	✓	✓	✓	0.95

5.10 Inferences Drawn from the Results

Four principal inferences emerge from the experimental findings. First, behavioural interaction data is a more reliable indicator of true modality preference than self-report: the 23.7 percentage point accuracy gap between static VARK profiling and hybrid dynamic profiling at Session 10 demonstrates that observed learning behaviour corrects the biases inherent in one-time questionnaire responses. Second, modality-aware content delivery confers the greatest benefit for abstract and procedural subjects — Digital Logic ($d = 1.24$) and Engineering Mathematics ($d = 1.10$) showed the largest effect sizes precisely because these subjects demand flexible mental representations that different modalities can supply in complementary ways. Third, personalisation sustains long-term engagement: the experimental group’s session duration continued to grow through Session 10, whereas the control group’s duration plateaued by Session 3, indicating that adaptive content forestalls the motivational decline associated with repetitive, uniform material. Fourth, accessibility-first architecture is not merely an ethical obligation but a performance multiplier: learners with accessibility needs achieved substantially higher learning gains (+24.6% vs. +9.3%) under accessibility-aware routing, confirming that designing for inclusion produces outcomes superior even to standard personalisation for this sub-group.

The results collectively validate all six research questions posed in Chapter 1, and the observed effect sizes compare favourably with those reported for existing adaptive learning platforms. The following chapter summarises the contributions of the NeuroLearn system, acknowledges the limitations of the current study, and outlines directions for future research including longitudinal retention studies, cross-domain validation, and integration of physiological sensing for richer learner modelling.

Chapter 6

Conclusion and Future Scope

CHAPTER 6

CONCLUSION AND FUTURE SCOPE

Every chapter should start with an introduction paragraph. This paragraph should brief about the flow of the chapter. This introduction can be limited within 4 to 5 sentences. The chapter heading should be appropriately modified (a sample heading is shown for this chapter). But don't start the introduction paragraph in the chapters like "This chapter deals with...". Instead you should bring in the highlights of the chapter in the introduction paragraph.

6.1 Conclusion

This chapter should not contain an introduction paragraph like other chapters. You can directly write conclusion of the work done under this section. Typically this section can have 3 to 4 paragraphs.

First paragraph should bring in the scenario of the project and every objective should be explained here.

Second paragraph should say how the objectives are implemented and achieved.

Last paragraph should draw the conclusions from each objective with quantitative results, performance improvement etc.

6.2 Future Scope

Briefly discuss the constraints and limitations of the project and state the possibilities of extending the work in future.

6.3 Learning Outcomes of the Project

- List the learning outcomes here
- List a minimum of 5 learning outcomes

Appendix A

Code

APPENDIX A

CODE

A.1 First Appendix

You can use `tcbllisting` for creating the code snippets. The following example illustrates how one can customize the `tcbllisting` to achieve the `tcl` script. Similarly, one can use it for other programming language listing, including HDL.

```
# Since our design has a clock with name clk,
## specify that name under [get_port ]
create_clock -period 40 -waveform {0 20} [get_ports clk]

# Setting a 'delay' on the clock:
set_clock_latency 0.3 clk

# Setting up constraints on your I/P and O/P pins
set_input_delay 2.0 -clock clk [all_inputs]
set_output_delay 1.65 -clock clk [all_outputs]

# Set realistic 'loads' on each output pin
set_load 0.1 [all_outputs]

# Set 'maximum' fanin and fan-out for the input and output pins
set_max_fanout 1 [all_inputs]
set_fanout_load 8 [all_outputs]
```